


# Random Acts of Pizza

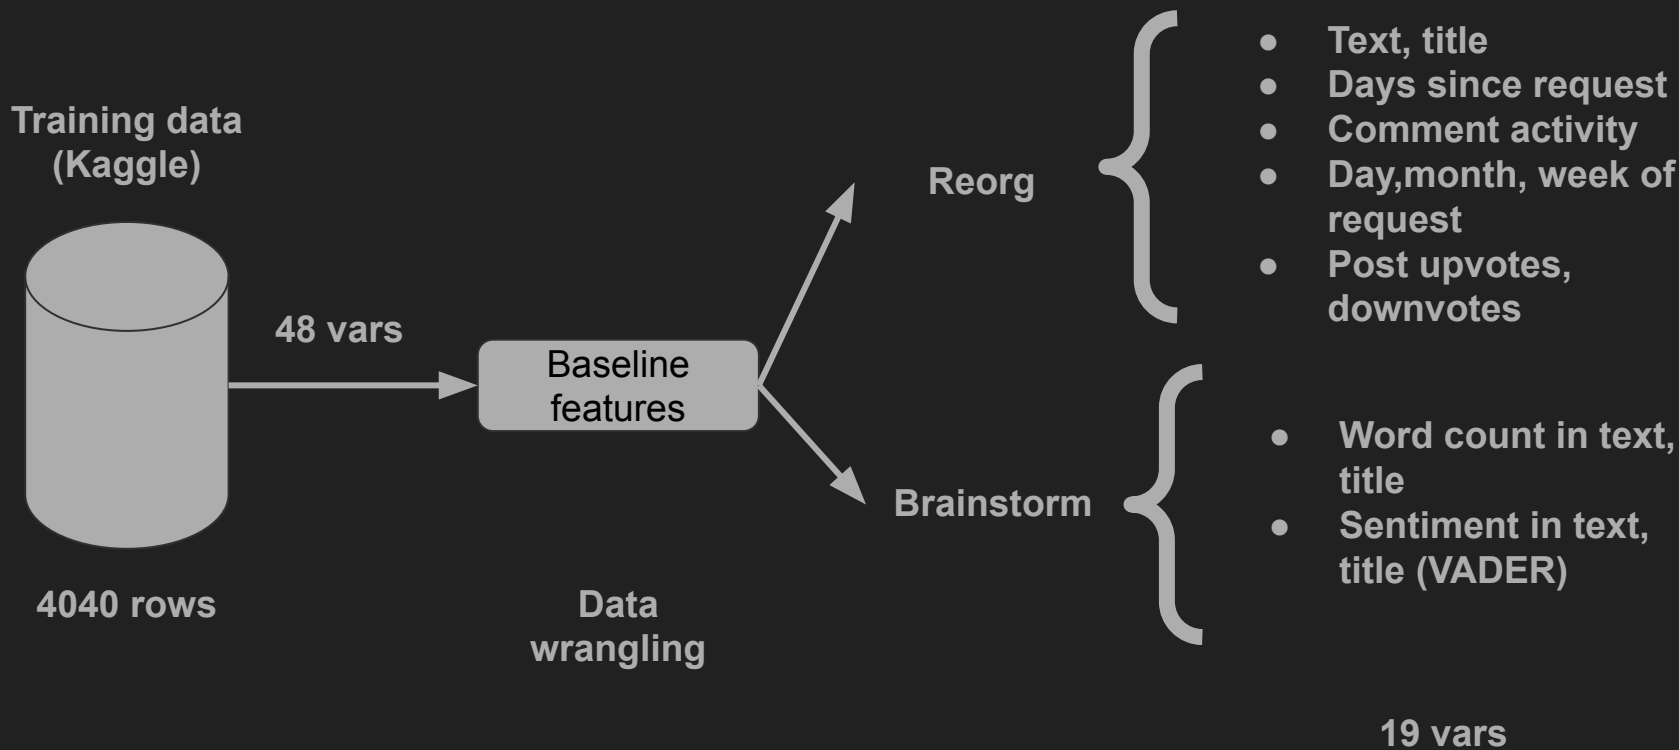
Kanika Mahajan, Carolina Arriaga  
W207 Applied Machine Learning

# Problem description

In the Random Acts of Pizza challenge (Kaggle competition) we are asked to **predict** if a message posted on Reddit group RAOP will get a pizza or not.



# Data description



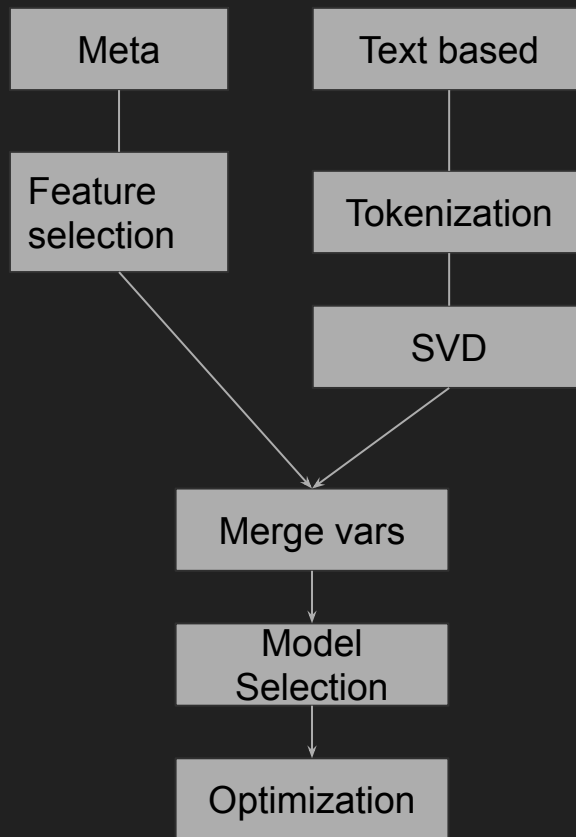
# Challenges

- The initial data was unbalanced.
  - 24% of the posts received pizza
  - Classifiers might be biased towards predicting “False”.
- We had a small data set (4,040 rows)
  - Test data not available (Kaggle competition)
  - Split in three ways -> even smaller dataset (75-15-10)
- Team objective:  
get accuracy > 90%



# Our Approach

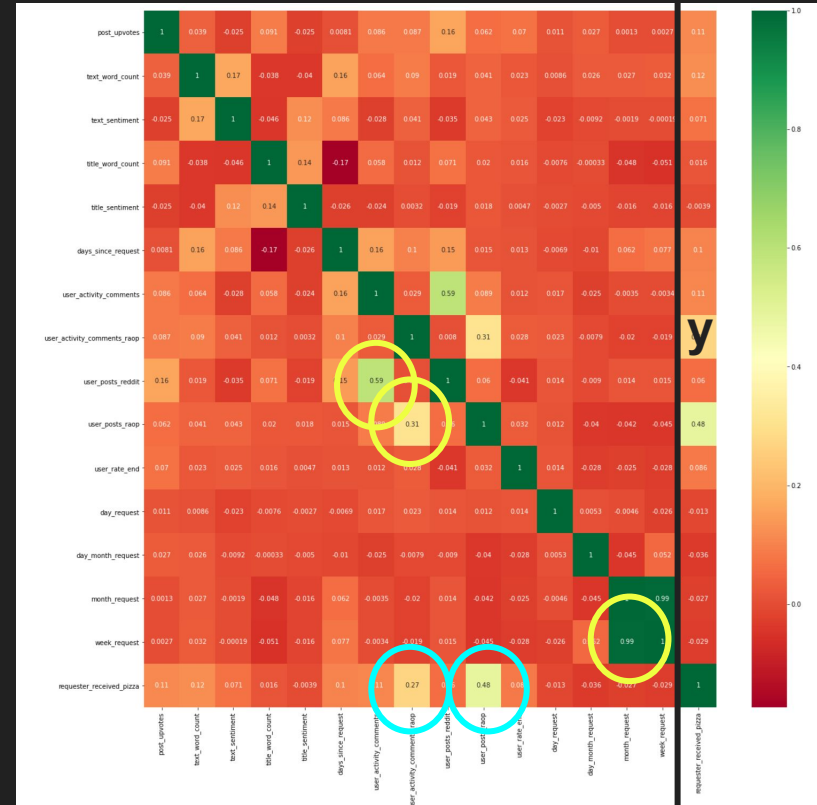
- Feature engineering
  - Meta - related to user or our understanding of language (e.g. post upvotes, activity)
  - Text based features - word counts, sentiment, TFIDF
  - Matrix correlation
  - Dimensionality reduction - SVD
- Model selection
  - Logistic regression
  - Random Forests (RF, Ada, XGB)
  - Multilayer Neural Network
  - Dense Neural Network
- Hyperparameter optimization



# Feature Selection - Meta

Finding features that are independent from each other to avoid multicollinearity.

- Highly correlated covariates:
  - Week - Month
  - Comments - Posts on Reddit
  - Comments raop - Posts on raop
- Removed week of request, posts on reddit, comments on raop.
- High correlation between output variable posts on raop (blue) - good variables.



Correlation matrix

# Feature engineering - TFIDF + SVD

We tried several approaches:


- Only tf-idf, word count (using token pattern)
- Better processor:
  - Removes special characters
  - Lemmatization and stemming
  - Remove stop words
- Used ngrams
  - Range (1-4)
- SVD

Tokenization

Up to 300k  
ngrams

SVD

Tested with:  
200, 400, 600,  
1000, 1500



Model Eval

Optimize Baseline  
Logistic  
regression

Random forest  
(XGB, Ada, RF)

# Model Overview - Baseline

## Logistic Regression

- Multiple componentes = 1500
- L2 Regularization

## Trees - (DT, RF, Ada)

- Estimators=20
- Depth=4

## 7-NN

- Tried [1,3,5,7,9,11]

Accuracy (Dev set)

Model	Accuracy
Logit	84%
DT	75%
RF	82%
Ada	83%
7-NN	72%



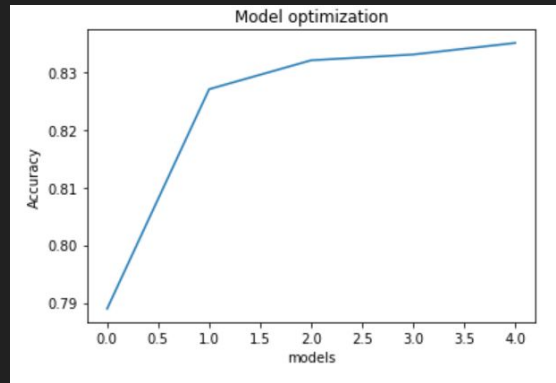
# Model Overview - Optimization

## Logistic Regression

- **SVD componentes** = [200, 400, 600, 1000, 1500]
- Multiple smoothing values  
**C**=[0.01,0.1,**1**,10,100]
- L2 Regularization

## Trees - (**XGB**, Ada, RF)

- Multiple estimators [15, 30, 80, 100, 120, **250**]
- Multiple depths [2,3,4,5,6,**12**]



Accuracy (Dev set)

Logit	XGB
84.5%	83%

1% improvement

# Model Overview - New models

## Multilayer Neural Network

- Used MLPClassifier with dimension (216,216,216) - our best result of components.

## Dense Neural Network

- Included 3 sigmoid layers
  - 2 of size 216
  - Third layer of size 1
  - Sigmoid

Accuracy (Dev set)

MLNN	DNN
82.17%	92.17%

+7% improvement

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 216)	46872
dense_1 (Dense)	(None, 216)	46872
dense_2 (Dense)	(None, 1)	217
Total params: 93,961		
Trainable params: 93,961		
Non-trainable params: 0		

# Model Overview - New model optimization

## Dense Neural Network

- Considered 1512 SVD (~60% variance)
- Stochastic Gradient Descent-  
batch size = 1
- Mini\_batches =  
[20,60,100,150]
- 2 Layers
- Units = [200 & 500] each layer
- Epochs = 100

Dev set	
Model	Accuracy
Stochastic GD	86%
Mini batch (20)	97%
Mini batch (60)	98%
Mini batch (150)	99%

+7% improvement

# DNN - Best model

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 200)	302600
dense_1 (Dense)	(None, 200)	40200
dropout (Dropout)	(None, 200)	0
dense_2 (Dense)	(None, 1)	201

Total params: 343,001

Trainable params: 343,001

Non-trainable params: 0

Accuracy (DNN): 0.994 - batch=150

Increased parameters

~ 3.6x

Increased accuracy to 99%

Added Noise (dropout=0.5)


Mini batch GD

# Optimized parameters - Best models

Model	Best hyperparameters	Accuracy
Logit	C = 1, L2, SVD=1512	85%
XGB	Depth=4 Estimators=120 SVD=1512	84%
DNN	SVD=1512 (input) Sequential L1= 200 comps & Sigmoid L2= 200 comps & Sigmoid Dropout = 0.5 L3= 1 comps & Sigmoid  Batch size=150 Epochs = 100  Optimizer=Adam	99%

# Results

We tested our optimized models with the validation data.



## Accuracy

Model	Dev	Test
Logit	86%	82%
XGB	84%	83%
DNN mini batch (150)	98%	99%

# What worked

- Removing covariates with high correlation.
- SVD components ~ hundreds.
- Logistic regression was slightly better than forests, also faster.
- Dense Neural Network predicted the best results.



# What didn't work

- Baseline model
  - KNN model, accuracy was much lower compared to RF and Logistic Regression.
- Multilayer Neural Network
  - Didn't do better than trees or logit.
  - Didn't explore further.
- Dense Neural Network
  - We tried different activation layers (relu, softmax) - sigmoid was the best





# Lessons learned

- We could have created an ensemble
  - Include multiple model predictions
- Some models were running well in a computer but due to computing power not so well in other.
  - We could have created a container
- Compiling takes a long time
  - Need to find ways to optimize computing time.

